

Planning Search Heuristic Analysis

Here a planning search agent is implemented to solve deterministic logistics planning problems for an Air Cargo transport system. We use a planning graph and automatic domain-independent heuristics with A* search and compare their results/performance against several uninformed non-heuristic search methods.

Three planning problems were given in the Air Cargo domain that use the same action schema:

Action(Load(c, p, a),

PRECOND: $At(c, a) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $\neg At(c, a) \wedge In(c, p)$

Action(Unload(c, p, a),

PRECOND: $In(c, p) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $At(c, a) \wedge \neg In(c, p)$

Action(Fly(p, from, to),

PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT: $\neg At(p, from) \wedge At(p, to)$

The three problems have the following initial states and goals:

Problem 1:

Init($At(C1, SFO) \wedge At(C2, JFK)$
 $\wedge At(P1, SFO) \wedge At(P2, JFK)$
 $\wedge Cargo(C1) \wedge Cargo(C2)$
 $\wedge Plane(P1) \wedge Plane(P2)$
 $\wedge Airport(JFK) \wedge Airport(SFO)$)
Goal($At(C1, JFK) \wedge At(C2, SFO)$)

Problem 2:

Init($At(C1, SFO) \wedge At(C2, JFK) \wedge At(C3, ATL)$
 $\wedge At(P1, SFO) \wedge At(P2, JFK) \wedge At(P3, ATL)$
 $\wedge Cargo(C1) \wedge Cargo(C2) \wedge Cargo(C3)$
 $\wedge Plane(P1) \wedge Plane(P2) \wedge Plane(P3)$
 $\wedge Airport(JFK) \wedge Airport(SFO) \wedge Airport(ATL)$)
Goal($At(C1, JFK) \wedge At(C2, SFO) \wedge At(C3, SFO)$)

Problem 3:

Init($At(C1, SFO) \wedge At(C2, JFK) \wedge At(C3, ATL) \wedge At(C4, ORD)$
 $\wedge At(P1, SFO) \wedge At(P2, JFK)$
 $\wedge Cargo(C1) \wedge Cargo(C2) \wedge Cargo(C3) \wedge Cargo(C4)$
 $\wedge Plane(P1) \wedge Plane(P2)$
 $\wedge Airport(JFK) \wedge Airport(SFO) \wedge Airport(ATL) \wedge Airport(ORD)$)
Goal($At(C1, JFK) \wedge At(C3, JFK) \wedge At(C2, SFO) \wedge At(C4, SFO)$)

The goals above can be reached using different plans, but the optimal plan lengths for problems 1,2, and 3 are 6, 9, and 12 actions, respectively. Below are sample plans with optimal length:

Problem 1:

Load(C1, P1, SFO)
Load(C2, P2, JFK)
Fly(P1, SFO, JFK)
Fly(P2, JFK, SFO)
Unload(C1, P1, JFK)
Unload(C2, P2, SFO)

Problem 2:

Load(C1, P1, SFO)
Load(C2, P2, JFK)
Load(C3, P3, ATL)
Fly(P1, SFO, JFK)
Fly(P2, JFK, SFO)
Fly(P3, ATL, SFO)
Unload(C3, P3, SFO)
Unload(C2, P2, SFO)
Unload(C1, P1, JFK)

Problem 3:

Load(C1, P1, SFO)
Load(C2, P2, JFK)
Fly(P1, SFO, ATL)
Load(C3, P1, ATL)
Fly(P2, JFK, ORD)
Load(C4, P2, ORD)
Fly(P1, ATL, JFK)
Fly(P2, ORD, SFO)
Unload(C4, P2, SFO)
Unload(C3, P1, JFK)
Unload(C2, P2, SFO)
Unload(C1, P1, JFK)

Uninformed Search Strategies Analysis

Search strategies that come under the heading of uninformed search have no additional information about states beyond that provided in the problem definition. All they can do is generate successors and distinguish a goal state from a non-goal state. In this section, we compare the performance of seven such strategies in terms of speed (execution time, measured in seconds), memory usage (measured in search node expansions) and optimality (Yes, if a solution of optimal length is found; No, otherwise). The number of goal tests and number of new nodes are not reported in the tables below since they do not change the results of our analysis below. Performance measures were collected using the following commands:

```
python run_search.py -p 1 -s 1 2 3 4 5 6 7 >> run_uninformed_search_results_p1.txt
python run_search.py -p 2 -s 1 3 5 7 >> run_uninformed_search_results_p2.txt
python run_search.py -p 3 -s 1 3 5 7 >> run_uninformed_search_results_p3.txt
```

For Problem 2, because their execution time exceeded 10 minutes, we cancelled data collection for Breadth First Tree Search, Depth Limited Search, and Recursive Best First Search (per Udacity instruction).

For the same reason, with Problem 3 we did not collect any data for Breadth First Tree Search, Depth Limited Search, Uniform Cost Search, and Recursive Best First Search.

Problem 1 Results

| Search Strategy | Optimal | Path Length | Execution Time(s) | Node Expansion |
|--------------------------------|---------|-------------|-------------------|----------------|
| Breadth First Search | Yes | 6 | 0.036 | 43 |
| Breadth First Tree Search | Yes | 6 | 1.06 | 1458 |
| Depth First Graph Search | No | 12 | 0.0087 | 12 |
| Depth Limited Search | No | 50 | 0.116 | 101 |
| Uniform Cost Search | Yes | 6 | 0.049 | 55 |
| Recursive Breadth First Search | Yes | 6 | 3.24 | 4229 |
| Greedy Best First Graph Search | Yes | 6 | 0.0086 | 7 |

Problem 2 Results

| Search Strategy | Optimal | Path Length | Execution Time(s) | Node Expansion |
|--------------------------------|---------|-------------|-------------------|----------------|
| Breadth First Search | Yes | 9 | 16.38 | 3401 |
| Breadth First Tree Search | - | - | - | - |
| Depth First Graph Search | No | 346 | 1.73 | 350 |
| Depth Limited Search | - | - | - | - |
| Uniform Cost Search | Yes | 9 | 52.4 | 4761 |
| Recursive Breadth First Search | - | - | - | - |
| Greedy Best First Graph Search | Yes | 9 | 3.67 | 550 |

Problem 3 Results

| Search Strategy | Optimal | Path Length | Execution Time(s) | Node Expansion |
|--------------------------------|---------|-------------|-------------------|----------------|
| Breadth First Search | Yes | 12 | 126.5 | 14491 |
| Breadth First Tree Search | - | - | - | - |
| Depth First Graph Search | No | 1878 | 26.38 | 1948 |
| Depth Limited Search | - | - | - | - |
| Uniform Cost Search | Yes | 12 | 527.1 | 17783 |
| Recursive Breadth First Search | - | - | - | - |
| Greedy Best First Graph Search | Yes | 12 | 92.85 | 4031 |

Analysis

With this 3-problem set, Breadth First Search and Uniform Cost Search are the only two uninformed search strategies that yield an optimal action plan under the 10mn time limit. When it comes to execution speed and memory usage, Depth First Graph Search is the fastest and uses the least memory. However, it does not generate an optimal action plan (problem 1: plan length of 12 instead of 6, problem 2: plan length of 346 instead of 9, problem 3: plan length of 1878 instead of 12).

If finding the optimal path length is critical, what strategy should we use?

Because it performs faster and uses less memory than Uniform Cost Search, Breadth First Search is the recommended search strategy. This isn't much of a surprise, as BFS is complete and optimal. Its only downside is memory usage, if the problem's branching factor is high

Which search strategy should we use, if having an optimal path length is not the primary criteria?

For problems 2 and 3, the Depth First Graph Search plan lengths are so much longer than the optimal path length that it wouldn't make sense to use this search strategy. Greedy Best First Graph Search is the best alternative. In problems 1 and 2, it manages to find the optimal path. In problem 3, it does not find the optimal path but the path length it generates is 22 instead of 10, which is much better than Depth First Graph Search (1878 path length!). Moreover, it still provides execution time savings and uses less memory than the best search strategy for an optimal solution (Breadth First Search).

Informed (Heuristic) Search Strategies Analysis

Informed search strategy — one that uses problem-specific knowledge beyond the definition of the problem itself — can find solutions more efficiently than can an uninformed strategy. In this section, we compare the performance of A* Search using three different heuristics. Here again, we evaluate these strategies in terms of speed, memory usage and optimality.

Performance measures were collected using the following commands:

```
python run_search.py -p 1 -s 8 9 10 >> run_informed_search_results_p1.txt
python run_search.py -p 2 -s 8 9 10 >> run_informed_search_results_p2.txt
python run_search.py -p 3 -s 8 9 >> run_informed_search_results_p3.txt
```

For Problems 2 and 3, because its execution time exceeded 10 minutes, we did not collect data for A* Search with Level Sum heuristic (as suggested by the Udacity).

For the same reason, with Problem 3 we did not collect any data for Breadth First Tree Search, Depth Limited Search, Uniform Cost Search, and Recursive Best First Search.

Problem 1 Results

| Search Strategy | Optimal | Path Length | Execution Time(s) | Node Expansion |
|--|----------------|--------------------|--------------------------|-----------------------|
| A* Search with h1 heuristic | Yes | 6 | 0.055 | 55 |
| A* Search with Ignore Precondition Heuristic | Yes | 6 | 0.049 | 41 |
| A* Search with Level Sum Heuristic | Yes | 6 | 3.45 | 11 |

Problem 2 Results

| Search Strategy | Optimal | Path Length | Execution Time(s) | Node Expansion |
|--|----------------|--------------------|--------------------------|-----------------------|
| A* Search with h1 heuristic | Yes | 9 | 54.29 | 4761 |
| A* Search with Ignore Precondition Heuristic | Yes | 9 | 17.47 | 1506 |
| A* Search with Level Sum Heuristic | Yes | 9 | 954.86 | 86 |

Problem 3 Results

| Search Strategy | Optimal | Path Length | Execution Time(s) | Node Expansion |
|--|----------------|--------------------|--------------------------|-----------------------|
| A* Search with h1 heuristic | Yes | 12 | 535.8 | 17783 |
| A* Search with Ignore Precondition Heuristic | Yes | 12 | 113.26 | 5081 |
| A* Search with Level Sum Heuristic | - | - | - | - |

Analysis

While all heuristics yield an optimal action plan, only the h1 and Ignore Preconditions heuristics return results within the 10mn max execution time set by the Udacity.

Which heuristic should we use? Of the two strategies mentioned above, A* Search with Ignore Preconditions heuristic is the fastest. If we let search run to completion on our machine, A* Search with Level Sum heuristic uses the least memory, but its execution time is much slower (16 mn for problem 2!).

Informed vs Uninformed Search Strategies

The search strategies that generate optimal plans are Breadth First Search, Uniform Cost Search, and A* Search with all three heuristics.

As we saw earlier, when it comes to execution speed and memory usage of uninformed search strategies, Depth First Graph Search is faster and uses less memory than Uniform Cost Search. As for informed search strategies, A* Search with Ignore Preconditions heuristic is the fastest and uses the least memory. So, really, the choice is between Depth First Graph Search and A* Search with Ignore Preconditions heuristic. Here we compare their results against our 3-problem set.

Problem 1 Results

| Search Strategy | Optimal | Path Length | Execution Time(s) | Node Expansion |
|--|----------------|--------------------|--------------------------|-----------------------|
| Breadth First Search | Yes | 6 | 0.036 | 43 |
| A* Search with Ignore Precondition Heuristic | Yes | 6 | 0.049 | 41 |

Problem 2 Results

| Search Strategy | Optimal | Path Length | Execution Time(s) | Node Expansion |
|--|----------------|--------------------|--------------------------|-----------------------|
| Breadth First Search | Yes | 9 | 16.38 | 3401 |
| A* Search with Ignore Precondition Heuristic | Yes | 9 | 17.47 | 1506 |

Problem 3 Results

| Search Strategy | Optimal | Path Length | Execution Time(s) | Node Expansion |
|--|----------------|--------------------|--------------------------|-----------------------|
| Breadth First Search | Yes | 12 | 126.5 | 14491 |
| A* Search with Ignore Precondition Heuristic | Yes | 12 | 113.26 | 5081 |

From the results above, because it is faster and uses less memory, A* Search with Ignore Preconditions heuristic would be the best choice overall for our Air Cargo problem

Conclusion

The results above clearly illustrate the benefits of using informed search strategies with custom heuristics over uninformed search techniques when searching for an optimal plan. The benefits are significant both in terms of speed and memory usage. Another, more subtle, benefit is that one can customize the trade-off between speed and memory usage by using different heuristics, which is simply not possible with uninformed search strategies.